

Owen-Ethan_905452983_palatics_Lab4

Ethan Owen

SID: 905452983

EE201A, Winter 2026, MSOL

Part I: Min Area

Initial Steps

I started by adding in all the files that we were given in the project as follows, and ensured they were referenced as needed in the script

- Gate-level netlist (`.v` file, generated by synthesis tool, e.g., Cadence Genus)
- Timing constraints (`.sdc` file, generated by synthesis tool)
- LEF for library/technology (`.lef` file, Nangate/FreePDK)
- Liberty timing specs for library (`.lib` file, Nangate)
- GDS for Nangate library cells (`.gds` file, to write final design GDS for tapeout)

UTIL Value Adjustments

I adjusted the UTIL value in the script and observed the resulting area that was generated from it.

Target Util	Core Area (um^2)	Chip Area (um^2)	Setup Slack (ns)	Hold Slack (ns)	DRC Violations	Result
0.90	383.838	1,005.480	0.527	0.017	0	PASS
0.99	349.258	948.024	0.518	0.017	0	PASS
0.991	349.258	948.024	0.518	0.017	0	PASS
0.992	(smaller)	--	--	--	781+	FAIL
0.995	(smaller)	--	--	--	781+	FAIL

I watched for improved cell utilizations as well, since a high cell utilization means that the design is crammed into a smaller footprint, and thus a smaller core area + chip area.

I landed on `0.991` being the best that I could go without introducing massive DRC violations.

Results

```
INITIAL_UTIL: 0.991
FINAL_UTIL: 0.98172
FINAL_SETUP_SLACK: 0.518
```

Part 2: No Timing Driven

Timing (hold, setup) and Power Reports

These were the end places where I saved results for power, setup time, and hold time, while running the

TCL script. These save points are consistent between the part 2 case with `timingDriven true` and `timingDriven false`.

- `before_placement`
- `after_placement`
- `after_opt_prects`
- `after_extractrc_postcts`
- `after_opt_postcts_hold`
- `after_extractrc_preroute`
- `after_global_route`
- `after_detail_route`
- `after_opt_postroute_setup`
- `after_opt_postroute_hold`
- `after_extractrc_postroute`
- `final_postroute`

Table of Information

This table was generated after running `lab4_part2` followed by `lab4_part2` with `timingDriven` set to false. The results are compared for hold & setup timing and power.

Step	Setup TD	Setup NoTD	Hold TD	Hold NoTD	Power TD	Power NoTD
Before Placement	0.754	0.754	0.013	0.013	0.24192193	0.24192193
After Placement	0.601	0.601	0.017	0.017	0.30237865	0.30237865
After Pre-CTS Opt	0.571	0.571	0.017	0.017	0.16421903	0.16421903
After Post-CTS RC Ext	0.582	0.582	0.017	0.017	0.16616036	0.16616036
After Post-CTS Hold Opt	0.582	0.582	0.017	0.017	0.16616036	0.16616036
After Pre-Route RC Ext	0.457	0.457	0.020	0.020	0.19107963	0.19107963
After Global Route	0.470	0.470	0.019	0.019	0.18469520	0.18469520
After Detail Route	0.470	0.470	0.019	0.019	0.18469520	0.18469520
After Post-Route Setup Opt	0.518	0.518	0.017	0.017	0.16604142	0.16604142
After Post-Route Hold Opt	0.518	0.518	0.017	0.017	0.16604142	0.16604142
Final Post-Route RC Ext	0.518	0.518	0.017	0.017	0.16604142	0.16604142

I noticed that the output of the final setup time was the same as that of part 1, which was a good sign that the results are consistent between the two steps.

What does `timingDriven` actually do?

Timing driven placement essentially makes timing a component of the placement objective. Normally, without this, we consider wirelength and congestion. But this option factors in timing as well. I think this mainly serves to improve critical timing paths in the place and after-placement optimization steps.

Results

I noticed *no* differences between `timingDriven false` and `timingDriven true`. There are a couple potential reasons for this that I see:

1. The design is pretty small
2. Both runs share the same core floorplan, utilization, libraries, constraints, so maybe `timingDriven` has a lesser effect due to that
3. Later optimization steps dominate the early ones with `timingDriven`
4. In my conclusions, I only compare worse case timing cases and power, there might be minor differences in the other results, but I doubt this.

Part 3: Power Structure

Adding Stripes in CLI

I used the CLI only and no GUI at all. To add the strips I used the following parameters:

```
set STRIPE_LAYER metal2
set STRIPE_DIRECTION vertical
set STRIPE_WIDTH_UM 0.21
set STRIPE_SPACING_UM 4.13
set STRIPE_OFFSET_LEFT_UM 2.0
set STRIPE_SET_DISTANCE_UM 8.26
```

These we extracted from the problem statement. The only outstanding figure here is the set distance, which ends up being double the stripe spacing.

Optimizations and UTIL

When I originally ran the script with no changes in UTIL values, I got huge DRC errors. In response I iteratively lowered the UTIL values from `0.991` to `0.918` where I finally passed DRC. Then I compared final values from here

Table of Comparison to Part 1

To compare against, I used my script from part 1 as the benchmark, and added in the metal strip and compared the results using a shell script.

Metric	Part 1	Part 3 Stripes	Delta (P3-P1)
Target Utilization	0.991	0.918	-0.073000
Final Utilization (%)	98.172	90.967	-7.205000
Worst Setup Slack (ns)	0.518	0.429	-0.089000
Worst Hold Slack (ns)	0.017	0.015	-0.002000
Total Power (mW)	---	0.19797320	---
Power Delta (%)	---	---	---
DRC Violations	0	0	0.000000
Core Area (um^2)	349.258	376.922	27.664000
Chip Area (um^2)	948.024	993.989	45.965000
Legality (Timing+DRC)	PASS	PASS	

Results

These results we accomplished with no DRC violations.

- Highest tested legal target utilization: 0.918
- Corresponding best final utilization (actual): 0.9067
- Worst setup slack: 0.429 ns
- Worst hold slack: 0.015 ns

By adding the power strips we worsened utilization, setup slack, and hold slack in comparison to the original TCL script that did not include the strips.